

[Home](#) » [Databricks](#) » [Data Analyst](#) »


Updated Databricks Certified Data Analyst Associate Dumps (V9.02) – Use the Latest Study Materials to Complete Your Exam

MAY 7, 2025

Updated Databricks Certified Data Analyst Associate Dumps (V9.02) – Use the Latest Study Materials to Complete Your Exam

Earning your Databricks Certified Data Analyst Associate certification is a strategic move to advance your data career. DumpsBase updated the Databricks Certified Data Analyst Associate dumps to V9.02, offering 59 practice exam questions and answers, which will be your latest study materials to pass the actual exam efficiently. These updated questions are expertly crafted to cover the entire exam objectives, all the answers have been verified, so with these Q&As, you can face the actual exam confidently. Prepare smartly and efficiently with Databricks Certified Data Analyst Associate dumps from DumpsBase. With verified Q&As, reliable support, and a 100% money-back guarantee, your [Databricks](#) Certified Data Analyst Associate certification success is within reach.

Check the **free dumps** of the **Databricks Certified Data Analyst Associate Dumps (V9.02):**



Latest Free Dumps

Continue to Read MB-240 Free Dumps (Part 2, Q41-Q80) of V13.02: Verify the Quality of the Microsoft Dynamics 365 Field Service Functional Consultant MB-240 Dumps July 7, 2025

400-007 Dumps (V19.03) Are the Most Current Version with 400 Q&As for Preparation: Check the 400-007 Free Dumps (Part 1, Q1-Q40) to Verify the Quality July 7, 2025

Valid EX370 Dumps for Your Red Hat Certified Specialist in OpenShift Data Foundation Exam Preparation: We have EX370 Free Dumps (Part 2,

Congratulations - you have completed this exam.

Your answers are shown below:

1. Which of the following layers of the medallion architecture is most commonly used by data analysts?

- None of these layers are used by data analysts
- Gold 
- All of these layers are used equally by data analysts
- Silver
- Bronze

Question was not answered

Explanation:

The gold layer of the medallion architecture contains data that is highly refined and aggregated, and powers analytics, machine learning, and production applications. Data analysts typically use the gold layer to access data that has been transformed into knowledge, rather than just information. The gold layer represents the final stage of data quality and optimization in the lakehouse.

Reference: What is the medallion lakehouse architecture?

2. A data analyst has recently joined a new team that uses Databricks SQL, but the analyst has never used Databricks before. The analyst wants to know where in Databricks SQL they can write and execute SQL queries.

On which of the following pages can the analyst write and execute SQL queries?

- Data page
- Dashboards page
- Queries page
- Alerts page
- SQL Editor page 

Question was not answered

Explanation:

The SQL Editor page is where the analyst can write and execute SQL queries in Databricks SQL. The SQL Editor page has a

Q41-Q80) of V8.02 Online July 7, 2025

PEGACPSA24V1 Free Dumps (Part 2, Q41-Q80): Offering 40 More Demo Questions to Help You Check the PEGACPSA24V1 Dumps (V8.02) July 7, 2025

Share ECBA Free Dumps (Part 2, Q41-Q60) of V10.02 to Verify the Quality: Get Ahead with Entry Certificate in Business Analysis (ECBA) Dumps July 7, 2025

Juniper JN0-253 Free Dumps (Part 3, Q81-Q120) of V8.02 Help You Check the Mist AI, Associate (JNCIA-MistAI) Exam Quality July 7, 2025

Microsoft MS-721 Dumps (V14.02) Accelerate Your Preparation: Continue to Check the MS-721 Free Dumps (Part 2, Q41-Q90) July 5, 2025

CAPM Dumps (V11.02) for Your Certified Associate in Project Management (CAPM) Exam Preparation: Check CAPM Free Dumps (Part 1, Q1-Q40) Today July 5, 2025

Share More Free Demos Online: Read Salesforce Process Automation Free Dumps (Part 2, Q41-Q65) of V13.02 Today July 5, 2025

CompTIA Linux+ XK0-005 Dumps Updated to V16.02 for Quick Preparation: Read XK0-005 Free Dumps (Part 1, Q1-Q40) Online July 5, 2025

Categories

(ISC)² (2)

query pane where the analyst can type or paste SQL statements, and a results pane where the analyst can view the query results in a table or a chart. The analyst can also browse data objects, edit multiple queries, execute a single query or multiple queries, terminate a query, save a query, download a query result, and more from the SQL Editor page.

Reference: Create a query in SQL editor

3. Which of the following describes how Databricks SQL should be used in relation to other business intelligence (BI) tools like Tableau, Power BI, and Looker?

- As an exact substitute with the same level of functionality
- As a substitute with less functionality
- As a complete replacement with additional functionality
- As a complementary tool for professional-grade presentations
- As a complementary tool for quick in-platform BI work 

Question was not answered

Explanation:

Databricks SQL is not meant to replace or substitute other BI tools, but rather to complement them by providing a fast and easy way to query, explore, and visualize data on the lakehouse using the built-in SQL editor, visualizations, and dashboards. Databricks SQL also integrates seamlessly with popular BI tools like Tableau, Power BI, and Looker, allowing analysts to use their preferred tools to access data through Databricks clusters and SQL warehouses. Databricks SQL offers low-code and no-code experiences, as well as optimized connectors and serverless compute, to enhance the productivity and performance of BI workloads on the lakehouse.

Reference: Databricks SQL, Connecting Applications and BI Tools to Databricks SQL, Databricks integrations

overview, Databricks SQL: Delivering a Production SQL Development Experience on the Lakehouse

4. Which of the following approaches can be used to connect Databricks to Fivetran for data ingestion?

- Use Workflows to establish a SQL warehouse (formerly known as a SQL endpoint) for Fivetran to interact with
- Use Delta Live Tables to establish a cluster for Fivetran to interact with

CAP Certification (1)

HCISPP Certification (1)

A10 Networks (1)

System Administration 4 (1)

AACE (4)

AACE Certification (1)

Professional Level
Certifications (3)

AACN (1)

CCRN (1)

AAFM (3)

Chartered Trust and Estate
Planner (1)

Chartered Wealth Manager (2)

AAPC (1)

Certified Professional Coder
(1)

ABA (2)

CRCM Certification (1)

ABPMP (1)

BPM Certification (1)

ACAMS (1)

ACE (1)

Personal Trainer (1)

ACFE (13)

Certified Fraud Examiner (13)

ACI (1)

ACI-Financial Markets
Association (1)

- Use Partner Connect's automated workflow to establish a cluster for Fivetran to interact with 
- Use Partner Connect's automated workflow to establish a SQL warehouse (formerly known as a SQL endpoint) for Fivetran to interact with
- Use Workflows to establish a cluster for Fivetran to interact with

Question was not answered

Explanation:

Partner Connect is a feature that allows you to easily connect your Databricks workspace to Fivetran and other ingestion partners using an automated workflow. You can select a SQL warehouse or a cluster as the destination for your data replication, and the connection details are sent to Fivetran. You can then choose from over 200 data sources that Fivetran supports and start ingesting data into Delta Lake.

Reference: Connect to Fivetran using Partner Connect, Use Databricks with Fivetran

5. Data professionals with varying titles use the Databricks SQL service as the primary touchpoint with the Databricks Lakehouse Platform. However, some users will use other services like Databricks Machine Learning or Databricks Data Science and Engineering.

Which of the following roles uses Databricks SQL as a secondary service while primarily using one of the other services?

- Business analyst
- SQL analyst
- Data engineer 
- Business intelligence analyst
- Data analyst

Question was not answered

Explanation:

Data engineers are primarily responsible for building, managing, and optimizing data pipelines and architectures. They use Databricks Data Science and Engineering service to perform tasks such as data ingestion, transformation, quality, and governance. Data engineers may use Databricks SQL as a

Acquia (1)

Drupal 8 Certification (1)

ACRP (1)

ACRP Certified Professional (1)

ACSM (2)

ACSM-CEP (2)

ACT (2)

ACT English Test (1)

ACT Mathematics Test (1)

Admission Test (2)

Graduate Management Admission Test (1)

Graduate Record Examinations (1)

Adobe (10)

Adobe Advertising (1)

Adobe Campaign Classic (2)

Adobe Commerce (2)

Adobe Experience Manager (1)

Adobe Real-Time CDP (1)

Adobe Workfront (3)

AFP (1)

Certified Treasury Professional (1)

AGA (1)

Government Financial Manager (1)

AHLEI (1)

secondary service to query, analyze, and visualize data from the lakehouse, but this is not their main focus.

Reference: Databricks SQL overview, Databricks Data Science and Engineering overview, Data engineering with Databricks

6. A data analyst has set up a SQL query to run every four hours on a SQL endpoint, but the SQL endpoint is taking too long to start up with each run.

Which of the following changes can the data analyst make to reduce the start-up time for the endpoint while managing costs?

- Reduce the SQL endpoint cluster size
- Increase the SQL endpoint cluster size
- Turn off the Auto stop feature
- Increase the minimum scaling value
- Use a Serverless SQL endpoint 

Question was not answered

Explanation:

A Serverless SQL endpoint is a type of SQL endpoint that does not require a dedicated cluster to run queries. Instead, it uses a shared pool of resources that can scale up and down automatically based on the demand. This means that a Serverless SQL endpoint can start up much faster than a SQL

endpoint that uses a cluster, and it can also save costs by only paying for the resources that are used. A Serverless SQL endpoint is suitable for ad-hoc queries and exploratory analysis, but it may not offer the same level of performance and isolation as a SQL endpoint that uses a cluster. Therefore, a data analyst should consider the trade-offs between speed, cost, and quality when choosing between a Serverless SQL endpoint and a SQL endpoint that uses a cluster.

Reference: Databricks SQL endpoints, Serverless SQL endpoints, SQL endpoint clusters

7. A data engineering team has created a Structured Streaming pipeline that processes data in micro-batches and populates gold-level tables. The microbatches are triggered every minute. A data analyst has created a dashboard based on this gold-level data. The project

Certified Hotel Administrator (CHA) (1)

AICPA (3)

AICPA Exam (1)

Public Accountant (1)

AIWMI (1)

Certified Credit Research Analyst (1)

Alcatel Lucent (7)

Alcatel-Lucent Certifications (7)

Alfresco (4)

ACSCA Certification (1)

Alfresco Certification (1)

APSCE Certification (1)

Process Services Administrator (1)

Alibaba Cloud (4)

Alibaba Developer (1)

AliCloud (14)

ACP Certification (1)

Alibaba Cloud Associate (ACA) (8)

Alibaba Cloud Certification Associate (1)

Alibaba Cloud Computing (4)

Amazon (59)

AWS Certification (14)

AWS Certified AI Practitioner (4)

AWS Certified Associate (20)

stakeholders want to see the results in the dashboard updated within one minute or less of new data becoming available within the gold-level tables.

Which of the following cautions should the data analyst share prior to setting up the dashboard to complete this task?

- The required compute resources could be costly 
The gold-level tables are not appropriately clean for
- business reporting
The streaming data is not an appropriate data source for a
- dashboard
- The streaming cluster is not fault tolerant
- The dashboard cannot be refreshed that quickly

Question was not answered

Explanation:

A Structured Streaming pipeline that processes data in micro-batches and populates gold-level tables every minute requires a high level of compute resources to handle the frequent data ingestion, processing, and writing. This could result in a significant cost for the organization, especially if the data volume and velocity are large. Therefore, the data analyst should share this caution with the project stakeholders before setting up the dashboard and evaluate the trade-offs between the desired refresh rate and the available budget. The other options are not valid cautions because:

B. The gold-level tables are assumed to be appropriately clean for business reporting, as they are the final output of the data engineering pipeline. If the data quality is not satisfactory, the issue should be addressed at the source or silver level, not at the gold level.

C. The streaming data is an appropriate data source for a dashboard, as it can provide near real-time insights and analytics for the business users. Structured Streaming supports various sources and sinks for streaming data, including Delta Lake, which can enable both batch and streaming queries on the same data.

D. The streaming cluster is fault tolerant, as Structured Streaming provides end-to-end exactly-once fault-tolerance guarantees through checkpointing and write-ahead logs. If a query fails, it can be restarted from the last checkpoint and resume processing.

AWS Certified Data Analytics – Specialty (2)

AWS Certified Database – Specialty (1)

AWS Certified Developer – Associate (1)

AWS Certified Professional (5)

AWS Certified Security – Specialty (1)

AWS Certified Specialty (6)

AWS Certified SysOps Administrator – Associate (1)

Data Engineer Associate (2)

DevOps-Engineer (1)

American College (1)

American College Certification (1)

AMPP (1)

Coatings Inspector Program (1)

Android (4)

Android Certified Application Developer (1)

Android certified Application Engineer (2)

Flutter Application Development (1)

APA (2)

AICP Certification (1)

American Planning Association (1)

APBM (1)

E. The dashboard can be refreshed within one minute or less of new data becoming available in the gold-level tables, as Structured Streaming can trigger micro-batches as fast as possible (every few seconds) and update the results incrementally. However, this may not be necessary or optimal for the business use case, as it could cause frequent changes in the dashboard and consume more resources.

Reference: Streaming on Databricks, Monitoring Structured Streaming queries on Databricks, A look at the new Structured Streaming UI in Apache Spark 3.0, Run your first Structured Streaming workload

8. Which of the following approaches can be used to ingest data directly from cloud-based object storage?

- Create an external table while specifying the DBFS storage path to FROM
- Create an external table while specifying the DBFS storage path to PATH
- It is not possible to directly ingest data from cloud-based object storage
- Create an external table while specifying the object storage path to FROM
- Create an external table while specifying the object storage path to LOCATION 

Question was not answered

Explanation:

External tables are tables that are defined in the Databricks metastore using the information stored in a cloud object storage location. External tables do not manage the data, but provide a schema and a table name to query the data. To create an external table, you can use the CREATE EXTERNAL TABLE statement and specify the object storage path to the LOCATION clause. For example, to create an external table named ext_table on a Parquet file stored in S3, you can use the following statement: SQL

```
CREATE EXTERNAL TABLE ext_table (  
  
col1 INT,  
  
col2 STRING  
  
)  
  
STORED AS PARQUET
```

CBM Certifications (1)

APEGS (1)

Professional Practice Exam (1)

API (4)

ICP Programs Certification (1)

Individual Certification
Programs (3)

APICS (8)

CLTD Certification (1)

CPIM Certification (3)

CSCP Certification (3)

Supply Chain Professional (1)

APM (4)

APM Project Fundamentals
Qualification (PFQ) (3)

Project professionals (1)

APMG-International (12)

Agile Business Analysis (2)

AgilePM (5)

Better Business Cases (2)

Change Management (3)

Appian (11)

Associate Developer (4)

Certified Analyst (1)

Lead Developer (3)

Senior Developer (3)

Apple (4)

Apple Certified iOS Technician
(ACiT) 2019 (1)

LOCATION 's3://bucket/path/file.parquet'

AI-generated code. Review and use carefully. More info on FAQ.

Reference: External tables

9. A data analyst wants to create a dashboard with three main sections: Development, Testing, and Production. They want all three sections on the same dashboard, but they want to clearly designate the sections using text on the dashboard.

Which of the following tools can the data analyst use to designate the Development, Testing, and Production sections using text?

- Separate endpoints for each section
- Separate queries for each section
- Markdown-based text boxes 
- Direct text written into the dashboard in editing mode
- Separate color palettes for each section

Question was not answered

Explanation:

Markdown-based text boxes are useful as labels on a dashboard. They allow the data analyst to add text to a dashboard using the %md magic command in a notebook cell and then select the dashboard

icon in the cell actions menu. The text can be formatted using markdown syntax and can include headings, lists, links, images, and more. The text boxes can be resized and moved around on the dashboard using the float layout option.

Reference: Dashboards in notebooks, How to add text to a dashboard in Databricks

10. A data analyst needs to use the Databricks Lakehouse Platform to quickly create SQL queries and data visualizations. It is a requirement that the compute resources in the platform can be made serverless, and it is expected that data visualizations can be placed within a dashboard.

Apple Certified IT Professional (ACITP) (1)

Apple Certified Support Professional (2)

APSE (1)

Employment Support Professional (1)

Arcitura (6)

Cloud Certified Professional (CCP) (5)

SOAArchitect (1)

ARDMS (1)

ARDMS SPI (1)

Arista (1)

Arista Certification (1)

Aruba (1)

Aruba ACMP (1)

ASCP Cert (1)

MLT(ASCP) (1)

ASIS (4)

Certified Protection Professional (3)

Professional Certified Investigator (1)

ASQ (9)

ASQ Certification (8)

Quality Improvement Associate (1)

ATD (1)

Talent Development (1)

Atlassian (17)

Which of the following Databricks Lakehouse Platform services/capabilities meets all of these requirements?

- Delta Lake
- Databricks Notebooks
- Tableau
- Databricks Machine Learning
- Databricks SQL 

Question was not answered

Explanation:

Databricks SQL is a serverless data warehouse on the Lakehouse that lets you run all of your SQL and BI applications at scale with your tools of choice, all at a fraction of the cost of traditional cloud data warehouses¹. Databricks SQL allows you to create SQL queries and data visualizations using the SQL Analytics UI or the Databricks SQL CLI². You can also place your data visualizations within a dashboard and share it with other users in your organization³. Databricks SQL is powered by Delta Lake, which provides reliability, performance, and governance for your data lake⁴.

Reference: Databricks SQL

Query data using SQL Analytics

Visualizations in Databricks notebooks

Delta Lake

11. A data analyst is attempting to drop a table `my_table`. The analyst wants to delete all table metadata and data.

They run the following command:

`DROP TABLE IF EXISTS my_table;`

While the object no longer appears when they run `SHOW TABLES`, the data files still exist.

Which of the following describes why the data files still exist and the metadata files were deleted?

- The table's data was larger than 10 GB
- The table did not have a location
- The table was external 
- The table's data was smaller than 10 GB

ACP-OA (2)

Jira Administrator (10)

Jira Project Administrator (4)

Jira service management (1)

Auldhouse (1)

MSPF Certification (1)

Autodesk (1)

Autodesk Maya (1)

Automation Anywhere (1)

RPA Certifications (1)

Avaya (95)

ACDS (7)

ACIS (27)

ACSS (38)

Administering Avaya Call Management System (ASAC-0021) (1)

APDS (3)

APSS (7)

Avaya Aura Communication Applications Technical Associate Support (ASTA-7230) (1)

Avaya Aura Communication Manager Administrator (ASAC-0038) (1)

Avaya Experience Portal with Proactive Outreach Manager Support (ACSS-3305) (1)

Avaya Experience Portal with Proactive Outreach Manager Technical Associate Support (ASTA-3305) (1)

- The table was managed

Question was not answered

Explanation:

An external table is a table that is defined in the metastore, but its data is stored outside of the Databricks environment, such as in S3, ADLS, or GCS. When an external table is dropped, only the metadata is deleted from the metastore, but the data files are not affected. This is different from a managed table, which is a table whose data is stored in the Databricks environment, and whose data files are deleted when the table is dropped. To delete the data files of an external table, the analyst needs to specify the PURGE option in the DROP TABLE command, or manually delete the files from the storage system.

Reference: DROP TABLE, Drop Delta table features, Best practices for dropping a managed Delta Lake table

12. After running DESCRIBE EXTENDED accounts.customers;, the following was returned:

Name	accounts.customers
Location	dbfs:/stakeholders/customers
Provider	delta
Owner	root
Type	EXTERNAL

Now, a data analyst runs the following command:

DROP accounts.customers;

Which of the following describes the result of running this command?

- Running SELECT * FROM delta.
- `dbfs:/stakeholders/customers` results in an error.
- Running SELECT * FROM accounts.customers will return all rows in the table.
- All files with the .customers extension are deleted.
- The accounts.customers table is removed from the metastore, and the underlying data files are deleted.
- The accounts.customers table is removed from the metastore, but the underlying data files are untouched. 

Question was not answered

Explanation:

Avaya IP Office™ Platform Administrator (ASAC-0013) (1)

Avaya IP Office™ Platform Technical Associate Implement (ASTA-7720) (1)

Avaya IP Office™ Platform Technical Associate Support (ASTA-3000) (2)

Avaya Solutions Administrator Certificate (1)

Avaya Workforce Engagement Technical Associate Support (ASTA-3308) (1)

Avaya-Certification (1)

AXP On-Prem (formerly Avaya Aura CC Elite) Support (ACSS-7492) (1)

AXP On-Prem (formerly Avaya Aura CC Elite) Technical Associate Implement (ASTA-7392) (1)

AXP On-Prem Technical Associate Support (ASTA – 7492) (1)

Aviatrix (2)

ACE Certifications (2)

AVIXA (1)

Certified Technology Specialist (CTS) (1)

Axis (2)

Axis Certified Professional (2)

BACB (2)

BACB certification (2)

BCM Institute (1)

BCMS Audit Certification (1)

the accounts.customers table is an EXTERNAL table, which means that it is stored outside the default warehouse directory and is not managed by Databricks. Therefore, when you run the DROP command on this table, it only removes the metadata information from the metastore, but does not delete the actual data files from the file system. This means that you can still access the data using the location path (dbfs:/stakeholders/customers) or create another table pointing to the same location. However, if you try to query the table using its name (accounts.customers), you will get an error because the table no longer exists in the metastore.

Reference: DROP TABLE | Databricks on AWS, Best practices for dropping a managed Delta Lake table - Databricks

13. Which of the following should data analysts consider when working with personally identifiable information (PII) data?

- Organization-specific best practices for PII data
Legal requirements for the area in which the data was collected
- None of these considerations
Legal requirements for the area in which the analysis is being performed
- All of these considerations 

Question was not answered

Explanation:

Data analysts should consider all of these factors when working with PII data, as they may affect the data security, privacy, compliance, and quality. PII data is any information that can be used to identify a specific individual, such as name, address, phone number, email, social security number, etc. PII data may be subject to different legal and ethical obligations depending on the context and location of the data collection and analysis. For example, some countries or regions may have stricter data protection laws than others, such as the General Data Protection Regulation (GDPR) in the European Union. Data analysts should also follow the organization-specific best practices for PII data, such as encryption, anonymization, masking, access control, auditing, etc. These best practices can help prevent data breaches, unauthorized access, misuse, or loss of PII data.

Reference: How to Use Databricks to Encrypt and Protect PII Data

Automating Sensitive Data (PII/PHI) Detection

BCS (17)

Artificial intelligence (AI) (3)

BCS Practitioner (1)

Business Analysis
Certifications (9)

Information security and CCP
scheme certifications (1)

ISTQB Certified Tester (1)

PRINCE2 Foundation (1)

Software testing (1)

BGC (1)

CPEA Certification (1)

BICSI (5)

RCDD Certification (5)

Blockchain (6)

Blockchain Certifications (2)

Certified Blockchain Developer
(2)

Certified Blockchain Solution
Architect (2)

BluePrism (17)

Developer (4)

Developer Certification (2)

Installation Engineer (2)

Professional Developer (1)

ROM Architect (4)

Solution Designer (2)

Technical Architect (2)

Broadcom (27)

Broadcom Technical Specialist

14. Delta Lake stores table data as a series of data files, but it also stores a lot of other information.

Which of the following is stored alongside data files when using Delta Lake?

- None of these
- Table metadata, data summary visualizations, and owner account information
- Table metadata 
- Data summary visualizations
- Owner account information

Question was not answered

Explanation:

Delta Lake is a storage layer that enhances data lakes with features like ACID transactions, schema enforcement, and time travel. While it stores table data as Parquet files, Delta Lake also keeps a transaction log (stored in the `_delta_log` directory) that contains detailed table metadata.

This metadata includes:

Table schema

Partitioning information

Data file paths

Transactional operations like inserts, updates, and deletes

Commit history and version control

This metadata is critical for supporting Delta Lake's advanced capabilities such as time travel and efficient query execution. Delta Lake does not store data summary visualizations or owner account information directly alongside the data files.

Reference: Delta Lake Table Features - Databricks Documentation

15. Which of the following is an advantage of using a Delta Lake-based data lakehouse over common data lake solutions?

(1)

Endpoint Security (2)

Symantec Certified Specialist (SCS) (1)

VCP-NV 2024 (2)

VCP-VCF Architect (4)

VMware vSAN Specialist (v2) (1)

C++ Institute (4)

C++ Certified Associate Programmer (2)

C++ Certified Professional Programmer (2)

CA (1)

CA Unified Infrastructure Management (1)

Certinia (2)

PSA Certified (2)

CertiProf (3)

DevOps Essentials Professional Certificate (1)

Lead Cybersecurity Professional Certificate (1)

Remote Work and Virtual Collaboration Certificate (1)

CertNexus (3)

Certified AI Practitioner (2)

Certified IoT Security Practitioner (1)

CFA Institute (6)

ESG Investing Certificate (2)

Level 1 CFA Exam (1)

- ACID transactions 
- Flexible schemas
- Data deletion
- Scalable storage
- Open-source formats

Question was not answered

Explanation:

A Delta Lake-based data lakehouse is a data platform architecture that combines the scalability and flexibility of a data lake with the reliability and performance of a data warehouse. One of the key advantages of using a Delta Lake-based data lakehouse over common data lake solutions is that it supports ACID transactions, which ensure data integrity and consistency. ACID transactions enable concurrent reads and writes, schema enforcement and evolution, data versioning and rollback, and data quality checks. These features are not available in traditional data lakes, which rely on file-based storage systems that do not support transactions.

Reference: Delta Lake: Lakehouse, warehouse, advantages | Definition Synapse C Data Lake vs. Delta Lake vs. Data Lakehouse Data Lake vs. Delta Lake - A Detailed Comparison

Building a Data Lakehouse with Delta Lake Architecture: A Comprehensive Guide

16. Which of the following benefits of using Databricks SQL is provided by Data Explorer?

- It can be used to run UPDATE queries to update any tables in a database.
- It can be used to view metadata and data, as well as view/change permissions. 
- It can be used to produce dashboards that allow data exploration.
- It can be used to make visualizations that can be shared with stakeholders.
- It can be used to connect to third party BI tools.

Question was not answered

Explanation:

Data Explorer is a user interface that allows you to discover and manage data, schemas, tables, models, and permissions in Databricks SQL. You can use Data Explorer to view schema

Level 2 CFA Exam (1)

Level 3 CFA Exam (2)

Check Point (47)

CCCS (3)

CCES (2)

CCME (4)

CCPA (1)

CCSA (9)

CCSE (15)

CCSM (2)

CCTA (1)

CCTE (7)

CCVS (1)

CTPS (1)

SandBlast (1)

CII (2)

CII credit level: Certificate (1)

CII credit level: Diploma (1)

CIMA (18)

Certificate in Business Accounting (6)

CIMA certification (3)

CIMA CGMA Professional Qualification (1)

Professional Qualification (9)

CIPS (37)

CIPS Certifications (1)

CIPS Level 3 Advanced Certificate in Procurement and Supply Operations (8)

details, preview sample data, and see table and model details and properties. Administrators can view and change owners, and admins and data object owners can grant and revoke

permissions¹.

Reference: Discover and manage data using Data Explorer

17. The stakeholders.customers table has 15 columns and 3,000 rows of data.

The following command is run:

```
CREATE TEMP VIEW stakeholders.eur_customers AS
  SELECT * FROM stakeholders.customers
  WHERE continent = 'eur';
```

After running SELECT * FROM stakeholders.eur_customers, 15 rows are returned. After the command executes completely, the user logs out of Databricks.

After logging back in two days later, what is the status of the stakeholders.eur_customers view?

- The view remains available and SELECT * FROM stakeholders.eur_customers will execute correctly. 
- The view has been dropped.

The view is not available in the metastore, but the underlying data can be accessed with SELECT * FROM delta.

- `stakeholders.eur\_customers`.

The view remains available but attempting to SELECT from it results in an empty result set because data in views are

- automatically deleted after logging out.
- The view has been converted into a table.

Question was not answered

Explanation:

In Databricks, a view is a saved SQL query definition that references existing tables or other views.

Once created, a view remains persisted in the metastore (such as Unity Catalog or Hive Metastore)

until it is explicitly dropped.

Key points:

CIPS Level 4 Diploma in Procurement and Supply (18)

CIPS Level 5 Advanced Diploma in Procurement and Supply (4)

CIPS Level 6 Professional Diploma in Procurement and Supply (6)

Cisco (651)

Account Manager (3)

Additional Online Exams for Validating Knowledge (2)

Advanced Collaboration Architecture Specialization (4)

Advanced Enterprise Networks Architecture Specialization (2)

Advanced Routing and Switching (1)

Advanced Unified Contact Center Enterprise Specialization (1)

Advanced Video Specialization (1)

AppDynamics Certified Associate (4)

AppDynamics Certified Implementation (2)

CCDA (4)

CCDE (12)

CCDP (10)

CCE (1)

CCEA (4)

CCECE (4)

CCENT (1)

Views do not store data themselves but reference data from underlying tables.

Logging out or being inactive does not delete or alter views.

Unless a user or admin explicitly drops the view or the underlying data/table is deleted, the view continues to function as expected.

Therefore, after logging back in—even days later—a user can still run `SELECT * FROM stakeholders.eur_customers`, and it will return the same data (provided the underlying table hasn't changed).

Reference: Views - Databricks Documentation

18. A data analyst created and is the owner of the managed table `my_table`. They now want to change ownership of the table to a single other user using Data Explorer.

Which of the following approaches can the analyst use to complete the task?

A. Edit the Owner field in the table page by removing their own account

B. Edit the Owner field in the table page by selecting All Users

C. Edit the Owner field in the table page by selecting the new owner's account

D. Edit the Owner field in the table page by selecting the Admins group

E. Edit the Owner field in the table page by removing all access

•

Explanation:

The Owner field in the table page shows the current owner of the table and allows the owner to change it to another user or group. To change the ownership of the table, the owner can click on the Owner field and select the new owner from the drop-down list. This will transfer the ownership of the table to the selected user or group and remove the previous owner from the list of table access control entries¹.

CCIE (15)

CCIE Collaboration (2)

CCIE Data Center (11)

CCIE Security (6)

CCIE Service Provider (6)

CCIE Wireless (2)

CCNA (31)

CCNA Cloud (4)

CCNA Collaboration (9)

CCNA Cyber Ops (8)

CCNA Data Center (7)

CCNA Routing and Switching (10)

CCNA Security (5)

CCNA SP (2)

CCNA Wireless (3)

CCNP (20)

CCNP Cloud (6)

CCNP Collaboration (56)

CCNP Data Center (67)

CCNP Enterprise (89)

CCNP Security (66)

CCNP Service Provider (24)

CCNP Wireless (13)

CCT Collaboration (1)

CCT Data Center (2)

CCT Field Technician (1)

CCT Routing & Switching (4)

The other options are incorrect because:

- A. Removing the owner's account from the Owner field will not change the ownership of the table, but will make the table ownerless².
- B. Selecting All Users from the Owner field will not change the ownership of the table, but will grant all users access to the table³.
- D. Selecting the Admins group from the Owner field will not change the ownership of the table, but will grant the Admins group access to the table³.
- E. Removing all access from the Owner field will not change the ownership of the table, but will revoke all access to the table⁴.

Reference: 1: Change table ownership

2: Ownerless tables

3: Table access control

4: Revoke access to a table

19. A data analyst has a managed table `table_name` in database `database_name`. They would now like to remove the table from the database and all of the data files associated with the table. The rest of the tables in the database must continue to exist.

Which of the following commands can the analyst use to complete the task without producing an error?

- `DROP DATABASE database_name;`
- `DROP TABLE database_name.table_name;` ✓
- `DELETE TABLE database_name.table_name;`
- `DELETE TABLE table_name FROM database_name;`
- `DROP TABLE table_name FROM database_name;`

Question was not answered

Explanation:

The `DROP TABLE` command removes a table from the metastore and deletes the associated data files. The syntax for this command is `DROP TABLE [IF EXISTS] [database_name.]table_name;`. The optional `IF EXISTS` clause prevents an error if the table does not exist. The optional

Cisco Business Architecture Analyst (5)

Cisco Business Architecture Practitioner (1)

Cisco Business Value Practitioner (1)

Cisco Business Value Specialist (1)

Cisco Certified CyberOps Associate (8)

Cisco Certified CyberOps Professional (4)

Cisco Certified DevNet Associate (7)

Cisco Certified DevNet Professional (13)

Cisco Certified Specialist (5)

Cisco Certified Support Technician (CCST) (4)

Cisco Channel Partner Data Center (1)

Cisco cloud (1)

Cisco Collaboration SaaS Authorization (1)

Cisco Collaboration Servers and Appliances Role (3)

Cisco Contact Center Enterprise (2)

Cisco Customer Success Manager Specialist (2)

Cisco Environmental Sustainability Specialization (2)

Cisco Meraki Solutions Specialist (3)

Cisco Network Programmability Design and Implementation

database_name. prefix specifies the database where the table resides. If not specified, the current database is used. Therefore, the correct command to remove the table table_name from the database database_name and all of the data files associated with it is DROP TABLE database_name.table_name;. The other commands are either invalid syntax or would produce undesired results.

Reference: Databricks - DROP TABLE

20. A data analyst runs the following command:

SELECT age, country

FROM my_table

WHERE age >= 75 AND country = 'canada';

Which of the following tables represents the output of the above command?

A)

age	country
80	canada
NULL	canada
90	NULL

B)

age	country
80	NULL
75	NULL
90	NULL

C)

id	age	country
900	80	canada
901	75	canada
902	90	canada

D)

Specialist (1)

Cisco Online Exams (1)

Cisco Renewals Manager (6)

Cisco Small and Medium
Business Engineer
Specialization (1)

Cisco Small and Medium
Business Sales Specialization
(3)

Cisco Specialist (2)

Cisco Specialist Certifications
(42)

Cisco Unified Contact Center
Enterprise Specialist (3)

CyberOps Associate (7)

ENUAE (2)

Express Collaboration
Specialization (6)

Express Security Specialization
(2)

Express Specialization (3)

Express Specialization Small
Business (3)

Other Cisco Certifications (9)

PreSales Engineer (1)

Systems Engineer (2)

CISI (1)

CISI level 3 Certificate in
Wealth & Investment
Management (1)

Citrix (34)

CCA-AppDS (5)

CCA-N (3)

age	country
80	canada
14	canada
90	canada

E)

age	country
80	canada
75	canada
90	canada

- Option A
- Option B
- Option C
- Option D
- Option E ✓

Question was not answered

Explanation:

The SQL query provided is designed to filter out records from "my_table" where the age is 75 or above and the country is Canada. Since I can't view the content of the links provided directly, I need to rely on the image attached to this question for context. Based on that, Option E (the image attached) represents a table with columns "age" and "country", showing records where age is 75 or above and country is Canada.

Reference: The answer can be inferred from understanding SQL queries and their outputs as per Databricks documentation: Databricks SQL

« **Updated Databricks
Certified Data Engineer
Associate Dumps (V11.02)
for Effective Preparation:
Choose DumpsBase to
Achieve Success**

**Databricks Machine
Learning Associate Dumps
(V9.02) 2025 - Pass Your
Databricks Certified
Machine Learning** »

CCA-V (9)

CCE-AppDS (1)

CCE-N (1)

CCE-V (5)

CCP-AppDS (2)

CCP-N (4)

CCP-V (6)

CIW (8)

CIW User Interface Designer
(1)

CIW Web and Mobile Design
Series (4)

CIW Web Development Series
(2)

Web Security Series (1)

CloudBees (2)

Jenkins Engineer (2)

Cloudera (1)

Cloudera Certified Associate
(CCA) (1)

CMT Association (1)

CMT Program (1)

CNCF (3)

CKA Certification (2)

CompTIA (261)

A+ (53)

CASP (5)

CompTIA Advanced Security
Practitioner (2)

CompTIA CASP (16)