


[Home](#) » [Databricks](#) » [Databricks Certification](#) »

Updated Databricks Certified Data Engineer Professional Dumps (V10.02) – Help You Access the Ability to Use Databricks to Perform Advanced Data Engineering Tasks

OCTOBER 23, 2023

Updated Databricks Certified Data Engineer Professional Dumps (V10.02) – Help You Access the Ability to Use Databricks to Perform Advanced Data Engineering Tasks

Once you pass the Databricks Certified Data Engineer Professional exam, you will achieve the prestigious Databricks Certified Data Engineer Professional certification. This certification is a testament to your proficiency in using Databricks to perform advanced data engineering tasks. DumpsBase has recently updated its Databricks Certified Data Engineer Professional dumps to V10.02, providing you with the perfect resource to ensure your success. With 337 practice exam questions and answers, DumpsBase guarantees that you will pass the Databricks Certified Data Engineer Professional exam smoothly. DumpsBase offers a 100% guarantee that you will pass the Databricks Certified Data Engineer Professional exam with the most updated exam dumps. Many Databricks exam students have successfully passed their Databricks Certified Professional



Latest Free Dumps

Continue to Read MB-240 Free Dumps (Part 2, Q41-Q80) of V13.02: Verify the Quality of the Microsoft Dynamics 365 Field Service Functional Consultant MB-240 Dumps July 7, 2025

400-007 Dumps (V19.03) Are the Most Current Version with 400 Q&As for Preparation: Check the 400-007 Free Dumps (Part 1, Q1-Q40) to Verify the Quality July 7, 2025

Valid EX370 Dumps for Your Red Hat Certified Specialist in OpenShift Data Foundation Exam Preparation: We have EX370 Free Dumps (Part 2,

Free Databricks Certified Data Engineer Professional Dumps Demos

Congratulations - you have completed this exam.

Your answers are shown below:

1. You were asked to setup a new all-purpose cluster, but the cluster is unable to start which of the following steps do you need to take to identify the root cause of the issue and the reason why the cluster was unable to start?

- Check the cluster driver logs
- Check the cluster event logs 
- (Correct)
- Workspace logs
- Storage account
- Data plane

Question was not answered

Explanation:

Cluster event logs are very useful, to identify issues pertaining to cluster availability. Cluster may not start due to resource limitations or issues with the cloud providers.

Some of the common issues include a subnet for compute VM reaching its limits or exceeding the subscription or cloud account CPU quota limit.

Here is an example where the cluster did not start due to subscription reaching the quota limit on a certain type of cpu cores for a VM type.



Q41-Q80) of V8.02 Online July 7, 2025

PEGACPSA24V1 Free Dumps (Part 2, Q41-Q80): Offering 40 More Demo Questions to Help You Check the PEGACPSA24V1 Dumps (V8.02) July 7, 2025

Share ECBA Free Dumps (Part 2, Q41-Q60) of V10.02 to Verify the Quality: Get Ahead with Entry Certificate in Business Analysis (ECBA) Dumps July 7, 2025

Juniper JN0-253 Free Dumps (Part 3, Q81-Q120) of V8.02 Help You Check the Mist AI, Associate (JNCIA-MistAI) Exam Quality July 7, 2025

Microsoft MS-721 Dumps (V14.02) Accelerate Your Preparation: Continue to Check the MS-721 Free Dumps (Part 2, Q41-Q90) July 5, 2025

CAPM Dumps (V11.02) for Your Certified Associate in Project Management (CAPM) Exam Preparation: Check CAPM Free Dumps (Part 1, Q1-Q40) Today July 5, 2025

Share More Free Demos Online: Read Salesforce Process Automation Free Dumps (Part 2, Q41-Q65) of V13.02 Today July 5, 2025

CompTIA Linux+ XK0-005 Dumps Updated to V16.02 for Quick Preparation: Read XK0-005 Free Dumps (Part 1, Q1-Q40) Online July 5, 2025

Categories

(ISC)² (2)

Graphical user

interface, text, application, email

Description automatically generated

Click on event logs



Graphical user

interface, text, application, email

Description automatically generated

Click on the message to see the detailed error message on why the cluster did not start.



Graphical user

interface, text, application, email

Description automatically generated

2. A SQL Dashboard was built for the supply chain team to monitor the inventory and product orders, but all of the timestamps displayed on the dashboards are showing in UTC format, so they requested to change the time zone to the location of New York.

How would you approach resolving this issue?

- Move the workspace from Central US zone to East US Zone
- Change the timestamp on the delta tables to America/New_York format
- Change the spark configuration of SQL endpoint to format the timestamp to America/New_York

CAP Certification (1)

HCISPP Certification (1)

A10 Networks (1)

System Administration 4 (1)

AACE (4)

AACE Certification (1)

Professional Level
Certifications (3)

AACN (1)

CCRN (1)

AAFM (3)

Chartered Trust and Estate
Planner (1)

Chartered Wealth Manager (2)

AAPC (1)

Certified Professional Coder
(1)

ABA (2)

CRCM Certification (1)

ABPMP (1)

BPM Certification (1)

ACAMS (1)

ACE (1)

Personal Trainer (1)

ACFE (13)

Certified Fraud Examiner (13)

ACI (1)

ACI-Financial Markets
Association (1)

- Under SQL Admin Console, set the SQL configuration 
- parameter time zone to America/New_York
 - Add SET Timezone = America/New_York on every of the SQL queries in the dashboard.

Question was not answered

Explanation:

The answer is, Under SQL Admin Console, set the SQL configuration parameter time zone to America/New_York

Here are steps you can take this to configure, so the entire dashboard is changed without

changing individual queries

Configure SQL parameters

To configure all warehouses with SQL parameters:

3. You are currently asked to work on building a data pipeline, you have noticed that you are currently working on a very large scale ETL many data dependencies, which of the following tools can be used to address this problem?

- AUTO LOADER
- JOBS and TASKS
- SQL Endpoints
- DELTA LIVE TABLES 
- STRUCTURED STREAMING with MULTI HOP

Question was not answered

Explanation:

The answer is, DELTA LIVE TABLES

DLT simplifies data dependencies by building DAG-based joins between live tables. Here is a view of how the dag looks with data dependencies without additional meta data,

4. When you drop a managed table using SQL syntax DROP TABLE table_name how does it impact metadata, history, and data stored in the table?

- Drops table from meta store, drops metadata, history, and 
- data in storage.

Acquia (1)

Drupal 8 Certification (1)

ACRP (1)

ACRP Certified Professional (1)

ACSM (2)

ACSM-CEP (2)

ACT (2)

ACT English Test (1)

ACT Mathematics Test (1)

Admission Test (2)

Graduate Management Admission Test (1)

Graduate Record Examinations (1)

Adobe (10)

Adobe Advertising (1)

Adobe Campaign Classic (2)

Adobe Commerce (2)

Adobe Experience Manager (1)

Adobe Real-Time CDP (1)

Adobe Workfront (3)

AFP (1)

Certified Treasury Professional (1)

AGA (1)

Government Financial Manager (1)

AHLEI (1)

- Drops table from meta store and data from storage but keeps metadata and history in storage
- Drops table from meta store, meta data and history but keeps the data in storage
- Drops table but keeps meta data, history and data in storage
- Drops table and history but keeps meta data and data in storage

Question was not answered

Explanation:

For a managed table, a drop command will drop everything from metastore and storage.

See the below image to understand the differences between dropping an external table.



Diagram

Description automatically generated

5. Which of the following approaches can the data engineer use to obtain a version-controllable configuration of the Job's schedule and configuration?

- They can link the Job to notebooks that are a part of a Databricks Repo.
- They can submit the Job once on a Job cluster.
- They can download the JSON equivalent of the job from the Job's page.
- They can submit the Job once on an all-purpose cluster. ✓
- They can download the XML description of the Job from the Job's page

Question was not answered

Certified Hotel Administrator (CHA) (1)

AICPA (3)

AICPA Exam (1)

Public Accountant (1)

AIWMI (1)

Certified Credit Research Analyst (1)

Alcatel Lucent (7)

Alcatel-Lucent Certifications (7)

Alfresco (4)

ACSCA Certification (1)

Alfresco Certification (1)

APSCE Certification (1)

Process Services Administrator (1)

Alibaba Cloud (4)

Alibaba Developer (1)

AliCloud (14)

ACP Certification (1)

Alibaba Cloud Associate (ACA) (8)

Alibaba Cloud Certification Associate (1)

Alibaba Cloud Computing (4)

Amazon (59)

AWS Certification (14)

AWS Certified AI Practitioner (4)

AWS Certified Associate (20)

6. What is the underlying technology that makes the Auto Loader work?

- Loader
- Delta Live Tables
- Structured Streaming ✓
- DataFrames
- Live DataFrames

Question was not answered

7. You are currently looking at a table that contains data from an e-commerce platform, each row contains a list of items(Item number) that were present in the cart, when the customer makes a change to the cart the entire information is saved as a separate list and appended to an existing list for the duration of the customer session, to identify all the items customer bought you have to make a unique list of items, you were asked to create a unique item's list that was added to the cart by the user, fill in the blanks of below query by choosing the appropriate higher-order function?

Note: See below sample data and expected output.

Schema: cartId INT, items Array<INT>

Sample data:

cartId	items
1	[[1,100,200,300], [1,250,300]]
2	[[10,150,200,300], [1,210,300],[350]]

Expected output:

cartId	items
1	[1,100,200,300,250]
2	[10,150,200,300,210,350]

Fill in the blanks:

Fill in the blanks:

SELECT cartId, **__**(items) **FROM** carts

- ARRAY_UNION, ARRAY_DISTINCT
- ARRAY_DISTINCT, ARRAY_UNION
- ARRAY_DISTINCT, FLATTEN ✓
- FLATTEN, ARRAY_DISTINCT
- ARRAY_DISTINCT, ARRAY_FLATTEN

AWS Certified Data Analytics – Specialty (2)

AWS Certified Database – Specialty (1)

AWS Certified Developer – Associate (1)

AWS Certified Professional (5)

AWS Certified Security – Specialty (1)

AWS Certified Specialty (6)

AWS Certified SysOps Administrator – Associate (1)

Data Engineer Associate (2)

DevOps-Engineer (1)

American College (1)

American College Certification (1)

AMPP (1)

Coatings Inspector Program (1)

Android (4)

Android Certified Application Developer (1)

Android certified Application Engineer (2)

Flutter Application Development (1)

APA (2)

AICP Certification (1)

American Planning Association (1)

APBM (1)

Question was not answered

Explanation:

FLATTEN -> Transforms an array of arrays into a single array.

ARRAY_DISTINCT -> The function returns an array of the same type as the input argument where all duplicate values have been removed.



Table

Description automatically generated

8. When building a DLT s pipeline you have two options to create a live tables, what is the main difference between CREATE STREAMING LIVE TABLE vs CREATE LIVE TABLE?

- CREATE STREAMING LIVE table is used in MULTI HOP
- Architecture
- CREATE LIVE TABLE is used when working with Streaming data sources and Incremental data
- CREATE STREAMING LIVE TABLE is used when working with Streaming data sources and Incremental data ✓
- There is no difference both are the same, CREATE STRAMING LIVE will be deprecated soon
- CREATE LIVE TABLE is used in DELTA LIVE TABLES, CREATE STREAMING LIVE can only used in Structured Streaming applications

Question was not answered

Explanation:

The answer is, CREATE STREAMING LIVE TABLE is used when working with Streaming data sources and Incremental data

9. You are tasked to set up a set notebook as a job for six departments and each department can run the task parallelly, the notebook takes an input parameter dept number to process the data by department, how do you go about to setup this up in job?

- Use a single notebook as task in the job and use dbutils.notebook.run to run each note-book with parameter in a different cell

CBM Certifications (1)

APEGS (1)

Professional Practice Exam (1)

API (4)

ICP Programs Certification (1)

Individual Certification Programs (3)

APICS (8)

CLTD Certification (1)

CPIM Certification (3)

CSCP Certification (3)

Supply Chain Professional (1)

APM (4)

APM Project Fundamentals Qualification (PFQ) (3)

Project professionals (1)

APMG-International (12)

Agile Business Analysis (2)

AgilePM (5)

Better Business Cases (2)

Change Management (3)

Appian (11)

Associate Developer (4)

Certified Analyst (1)

Lead Developer (3)

Senior Developer (3)

Apple (4)

Apple Certified iOS Technician (ACiT) 2019 (1)

A task in the job cannot take an input parameter, create six notebooks with hardcoded dept number and setup six tasks

- with linear dependency in the job

A task accepts key-value pair parameters, creates six tasks pass department number as parameter foreach task with no dependency in the job as they can all run in parallel. (Correct) ✓

A parameter can only be passed at the job level, create six jobs pass department number to each job with linear job dependency

A parameter can only be passed at the job level, create six jobs pass department number to each job with no job dependency

Question was not answered

Explanation:

Here is how you setup

Create a single job and six tasks with the same notebook and assign a different parameter for each task,



Graphical user interface, text, application, email

Description automatically generated

Apple Certified IT Professional (ACITP) (1)

Apple Certified Support Professional (2)

APSE (1)

Employment Support Professional (1)

Arcitura (6)

Cloud Certified Professional (CCP) (5)

SOA Architect (1)

ARDMS (1)

ARDMS SPI (1)

Arista (1)

Arista Certification (1)

Aruba (1)

Aruba ACMP (1)

ASCP Cert (1)

MLT(ASCP) (1)

ASIS (4)

Certified Protection Professional (3)

Professional Certified Investigator (1)

ASQ (9)

ASQ Certification (8)

Quality Improvement Associate (1)

ATD (1)

Talent Development (1)

Atlassian (17)

All tasks are added in a single job and can run parallel either using single shared cluster or with individual clusters.



Graphical user

interface, application, Teams

Description automatically generated

10. Which of the following commands can be used to run one notebook from another notebook?

- `notebook.utils.run("full notebook path")`
- `execute.utils.run("full notebook path")`
- `dbutils.notebook.run("full notebook path")` ✓
- only job clusters can run notebook
- `spark.notebook.run("full notebook path")`

Question was not answered

Explanation:

The answer is `dbutils.notebook.run("full notebook path")`

Here is the full command with additional options.

`run(path: String, timeout_seconds: int, arguments: Map): String`

11. You have configured AUTO LOADER to process incoming IOT data from cloud object storage every 15 mins, recently a change was made to the notebook code to update the processing logic but the team later realized that the notebook was failing for the last 24 hours, what steps team needs to take to reprocess the data that was not loaded after the notebook was corrected?

- Move the files that were not processed to another location and manually copy the files into the ingestion path to reprocess them
- Enable `back_fill = TRUE` to reprocess the data
- Delete the checkpoint folder and run the autoloader again
Autoloader automatically re-processes data that was not loaded ✓
- Manually re-load the data

Question was not answered

ACP-OA (2)

Jira Administrator (10)

Jira Project Administrator (4)

Jira service management (1)

Auldhouse (1)

MSPF Certification (1)

Autodesk (1)

Autodesk Maya (1)

Automation Anywhere (1)

RPA Certifications (1)

Avaya (95)

ACDS (7)

ACIS (27)

ACSS (38)

Administering Avaya Call Management System (ASAC-0021) (1)

APDS (3)

APSS (7)

Avaya Aura Communication Applications Technical Associate Support (ASTA-7230) (1)

Avaya Aura Communication Manager Administrator (ASAC-0038) (1)

Avaya Experience Portal with Proactive Outreach Manager Support (ACSS-3305) (1)

Avaya Experience Portal with Proactive Outreach Manager Technical Associate Support (ASTA-3305) (1)

Explanation:

The answer is,

Autoloader automatically re-processes data that was not loaded using the checkpoint.

12. John Smith is a newly joined team member in the Marketing team who currently has access read access to sales tables but does not have access to delete rows from the table, which of the following commands help you accomplish this?

- GRANT USAGE ON TABLE table_name TO
• john.smith@marketing.com
- GRANT DELETE ON TABLE table_name TO
• john.smith@marketing.com
- GRANT DELETE TO TABLE table_name ON
• john.smith@marketing.com
- GRANT MODIFY TO TABLE table_name ON
• john.smith@marketing.com
- GRANT MODIFY ON TABLE table_name TO
• john.smith@marketing.com



Question was not answered

Explanation:

The answer is GRANT MODIFY ON TABLE table_name TO john.smith@marketing.com, please note INSERT, UPDATE, and DELETE are combined into one role called MODIFY. Below are the list of privileges that can be granted to a user or a group, SELECT: gives read access to an object.

CREATE: gives the ability to create an object (for example, a table in a schema).

MODIFY: gives the ability to add, delete, and modify data to or from an object.

USAGE: does not give any abilities, but is an additional requirement to perform any action on a schema object.

READ_METADATA: gives the ability to view an object and its metadata.

CREATE_NAMED_FUNCTION: gives the ability to create a named UDF in an existing catalog or schema.

Avaya IP Office™ Platform
Administrator (ASAC-0013) (1)

Avaya IP Office™ Platform
Technical Associate Implement
(ASTA-7720) (1)

Avaya IP Office™ Platform
Technical Associate Support
(ASTA-3000) (2)

Avaya Solutions Administrator
Certificate (1)

Avaya Workforce Engagement
Technical Associate Support
(ASTA-3308) (1)

Avaya-Certification (1)

AXP On-Prem (formerly Avaya
Aura CC Elite) Support (ACSS-
7492) (1)

AXP On-Prem (formerly Avaya
Aura CC Elite) Technical
Associate Implement (ASTA-
7392) (1)

AXP On-Prem Technical
Associate Support (ASTA –
7492) (1)

Aviatrix (2)

ACE Certifications (2)

AVIXA (1)

Certified Technology Specialist
(CTS) (1)

Axis (2)

Axis Certified Professional (2)

BACB (2)

BACB certification (2)

BCM Institute (1)

BCMS Audit Certification (1)

MODIFY_CLASSPATH: gives the ability to add files to the Spark classpath.

ALL PRIVILEGES: gives all privileges (is translated into all the above privileges)

13. Which of the following SQL command can be used to insert or update or delete rows based on a condition to check if a row(s) exists?

- MERGE INTO table_name ✓
- COPY INTO table_name
- UPDATE table_name
- INSERT INTO OVERWRITE table_name
- INSERT IF EXISTS table_name

Question was not answered

Explanation:

here is the additional documentation for your review.

<https://docs.databricks.com/spark/latest/spark-sql/language-manual/delta-merge-into.html>

14. return x

check_input(1,3)

- 1
- 2
- 3 (Correct) ✓
- 4
- 5

Question was not answered

15. Which of the following is true, when building a Databricks SQL dashboard?

- A dashboard can only use results from one query
Only one visualization can be developed with one query result
- A dashboard can only connect to one schema/Database
More than one visualization can be developed using a single query result ✓
- A dashboard can only have one refresh schedule

Question was not answered

BCS (17)

Artificial intelligence (AI) (3)

BCS Practitioner (1)

Business Analysis
Certifications (9)

Information security and CCP
scheme certifications (1)

ISTQB Certified Tester (1)

PRINCE2 Foundation (1)

Software testing (1)

BGC (1)

CPEA Certification (1)

BICSI (5)

RCDD Certification (5)

Blockchain (6)

Blockchain Certifications (2)

Certified Blockchain Developer
(2)

Certified Blockchain Solution
Architect (2)

BluePrism (17)

Developer (4)

Developer Certification (2)

Installation Engineer (2)

Professional Developer (1)

ROM Architect (4)

Solution Designer (2)

Technical Architect (2)

Broadcom (27)

Broadcom Technical Specialist

Explanation:

the answer is, More than one visualization can be developed using a single query result. In the query editor pane + Add visualization tab can be used for many visualizations for a single query result.



Graphical user

interface, text, application

Description automatically generated

16. What is the main difference between the bronze layer and silver layer in a medallion architecture?

- Duplicates are removed in bronze, schema is applied in silver
- Silver may contain aggregated data
Bronze is raw copy of ingested data, silver contains data with production schema and optimized for ELT/ETL ✓
- throughput
- Bad data is filtered in Bronze, silver is a copy of bronze data

Question was not answered

Explanation:

Medallion Architecture C Databricks

Exam focus: Please review the below image and understand the role of each layer(bronze, silver, gold) in medallion architecture, you will see varying questions targeting each layer and its purpose.

Sorry I had to add the watermark some people in Udemy are copying my content.

(1)

Endpoint Security (2)

Symantec Certified Specialist (SCS) (1)

VCP-NV 2024 (2)

VCP-VCF Architect (4)

VMware vSAN Specialist (v2) (1)

C++ Institute (4)

C++ Certified Associate Programmer (2)

C++ Certified Professional Programmer (2)

CA (1)

CA Unified Infrastructure Management (1)

Certinia (2)

PSA Certified (2)

CertiProf (3)

DevOps Essentials Professional Certificate (1)

Lead Cybersecurity Professional Certificate (1)

Remote Work and Virtual Collaboration Certificate (1)

CertNexus (3)

Certified AI Practitioner (2)

Certified IoT Security Practitioner (1)

CFA Institute (6)

ESG Investing Certificate (2)

Level 1 CFA Exam (1)



A diagram of a house

Description automatically generated with low confidence

17. You noticed that a team member started using an all-purpose cluster to develop a notebook and used the same all-purpose cluster to set up a job that can run every 30 mins so they can update un-derlying tables which are used in a dashboard.

What would you recommend for reducing the overall cost of this approach?

- Reduce the size of the cluster
- Reduce the number of nodes and enable auto scale
- Enable auto termination after 30 mins
- Change the cluster all-purpose to job cluster when scheduling the job
- Change the cluster mode from all-purpose to single-mode



Question was not answered

Explanation:

While using an all-purpose cluster is ok during development but anytime you don't need to interact with a notebook, especially for a scheduled job it is less expensive to use a job cluster. Using an all-purpose cluster can be twice as expensive as a job cluster.

Please note: The compute cost you pay the cloud provider for the same cluster type and size be-tween an all-purpose cluster and job cluster is the same the only difference is the DBU cost.

The total cost of cluster = Total cost of VM compute (Azure or AWS or GCP) + Cost per DBU

The per DBU cost varies between all-purpose and Job Cluster

Level 2 CFA Exam (1)

Level 3 CFA Exam (2)

Check Point (47)

CCCS (3)

CCES (2)

CCME (4)

CCPA (1)

CCSA (9)

CCSE (15)

CCSM (2)

CCTA (1)

CCTE (7)

CCVS (1)

CTPS (1)

SandBlast (1)

CII (2)

CII credit level: Certificate (1)

CII credit level: Diploma (1)

CIMA (18)

Certificate in Business
Accounting (6)

CIMA certification (3)

CIMA CGMA Professional
Qualification (1)

Professional Qualification (9)

CIPS (37)

CIPS Certifications (1)

CIPS Level 3 Advanced
Certificate in Procurement and
Supply Operations (8)

Here is the recent cost estimate from AWS between Jobs Cluster and all-purpose Cluster, for jobs compute its \$0.15 cents per DBU v\$0.55 cents per DBU for all-purpose



Graphical user

interface

Description automatically generated

How do I check how much the DBU cost for my cluster?

When you click on an exister cluster or when you look at the cluster details you will see this in the top right corner



Graphical user

interface, text, application, email

Description automatically generated

CIPS Level 4 Diploma in
Procurement and Supply (18)

CIPS Level 5 Advanced
Diploma in Procurement and
Supply (4)

CIPS Level 6 Professional
Diploma in Procurement and
Supply (6)

Cisco (651)

Account Manager (3)

Additional Online Exams for
Validating Knowledge (2)

Advanced Collaboration
Architecture Specialization (4)

Advanced Enterprise Networks
Architecture Specialization (2)

Advanced Routing and
Switching (1)

Advanced Unified Contact
Center Enterprise
Specialization (1)

Advanced Video Specialization
(1)

AppDynamics Certified
Associate (4)

AppDynamics Certified
Implementation (2)

CCDA (4)

CCDE (12)

CCDP (10)

CCE (1)

CCEA (4)

CCECE (4)

CCENT (1)

18. Which of the following locations hosts the driver and worker nodes of a Databricks-managed cluster?

- Data plane ✓
- Control plane
- Databricks Filesystem
- JDBC data source
- Databricks web application

Question was not answered

Explanation:

The answer is Data Plane, which is where compute(all-purpose, Job Cluster, DLT) are stored this is generally a customer cloud account, there is one exception SQL Warehouses, currently there are 3 types of SQL Warehouse compute available(classic, pro, serverless), in classic and pro compute is located in customer cloud account but serverless computed is located in Databricks cloud account.



Diagram,

timeline

Description automatically generated

19. Which of the following data workloads will utilize a Bronze table as its destination?

- CCIE (15)
- CCIE Collaboration (2)
- CCIE Data Center (11)
- CCIE Security (6)
- CCIE Service Provider (6)
- CCIE Wireless (2)
- CCNA (31)
- CCNA Cloud (4)
- CCNA Collaboration (9)
- CCNA Cyber Ops (8)
- CCNA Data Center (7)
- CCNA Routing and Switching (10)
- CCNA Security (5)
- CCNA SP (2)
- CCNA Wireless (3)
- CCNP (20)
- CCNP Cloud (6)
- CCNP Collaboration (56)
- CCNP Data Center (67)
- CCNP Enterprise (89)
- CCNP Security (66)
- CCNP Service Provider (24)
- CCNP Wireless (13)
- CCT Collaboration (1)
- CCT Data Center (2)
- CCT Field Technician (1)
- CCT Routing & Switching (4)

- A job that aggregates cleaned data to create standard summary statistics
- A job that queries aggregated data to publish key insights into a dashboard
- A job that ingests raw data from a streaming source into the Lakehouse 
- A job that develops a feature set for a machine learning application
- A job that enriches data by parsing its timestamps into a human-readable format

Question was not answered

Explanation:

The answer is A job that ingests raw data from a streaming source into the Lakehouse.

The ingested data from the raw streaming data source like Kafka is first stored in the Bronze layer as first destination before it is further optimized and stored in Silver.

Medallion Architecture C Databricks

Bronze Layer:

20. What could be the expected output of query `SELECT COUNT (DISTINCT *) FROM user` on this table

userId	username	email
1	john.smith	john.smith@com
2	NULL	david.clear@com
3	kevin.smith	kevin.smith@com

- 3
- 2 (Correct) 
- 1
- 0
- NULL

Question was not answered

Explanation:

The answer is 2,

Count (DISTINCT *) removes rows with any column with a NULL value

Cisco Business Architecture Analyst (5)

Cisco Business Architecture Practitioner (1)

Cisco Business Value Practitioner (1)

Cisco Business Value Specialist (1)

Cisco Certified CyberOps Associate (8)

Cisco Certified CyberOps Professional (4)

Cisco Certified DevNet Associate (7)

Cisco Certified DevNet Professional (13)

Cisco Certified Specialist (5)

Cisco Certified Support Technician (CCST) (4)

Cisco Channel Partner Data Center (1)

Cisco cloud (1)

Cisco Collaboration SaaS Authorization (1)

Cisco Collaboration Servers and Appliances Role (3)

Cisco Contact Center Enterprise (2)

Cisco Customer Success Manager Specialist (2)

Cisco Environmental Sustainability Specialization (2)

Cisco Meraki Solutions Specialist (3)

Cisco Network Programmability Design and Implementation

21. A DELTA LIVE TABLE pipelines can be scheduled to run in two different modes, what are these two different modes?

- Triggered, Incremental
- Once, Continuous
- Triggered, Continuous 
- Once, Incremental
- Continuous, Incremental

Question was not answered

Explanation:

The answer is Triggered, Continuous

<https://docs.microsoft.com/en-us/azure/databricks/data-engineering/delta-live-tables/delta-live-tables-concepts#--continuous-and-triggered-pipelines>

- Triggered pipelines update each table with whatever data is currently available and then stop the cluster running the pipeline. Delta Live Tables automatically analyzes the dependencies between your tables and starts by computing those that read from external sources. Tables within the pipeline are updated after their dependent data sources have been updated.
- Continuous pipelines update tables continuously as input data changes. Once an update is started, it continues to run until manually stopped. Continuous pipelines require an always-running cluster but ensure that downstream consumers have the most up-to-date data.

22. Which of the following developer operations in CI/CD flow can be implemented in Databricks Re-pos?

- Merge when code is committed
- Pull request and review process
- Trigger Databricks Repos API to pull the latest version of code into production folder 
- Resolve merge conflicts
- Delete a branch

Question was not answered

Explanation:

Specialist (1)

Cisco Online Exams (1)

Cisco Renewals Manager (6)

Cisco Small and Medium Business Engineer Specialization (1)

Cisco Small and Medium Business Sales Specialization (3)

Cisco Specialist (2)

Cisco Specialist Certifications (42)

Cisco Unified Contact Center Enterprise Specialist (3)

CyberOps Associate (7)

ENUAE (2)

Express Collaboration Specialization (6)

Express Security Specialization (2)

Express Specialization (3)

Express Specialization Small Business (3)

Other Cisco Certifications (9)

PreSales Engineer (1)

Systems Engineer (2)

CISI (1)

CISI level 3 Certificate in Wealth & Investment Management (1)

Citrix (34)

CCA-AppDS (5)

CCA-N (3)

See the below diagram to understand the role Databricks Repos and Git provider plays when building a CI/CD workflow.

All the steps highlighted in yellow can be done Databricks Repo, all the steps highlighted in Gray are done in a git provider like Github or Azure DevOps



23. Identify one of the below statements that can query a delta table in PySpark Dataframe API

- Spark.read.mode("delta").table("table_name")
- Spark.read.table.delta("table_name")
- Spark.read.table("table_name") ✓
- Spark.read.format("delta").LoadTableAs("table_name")
- Spark.read.format("delta").TableAs("table_name")

Question was not answered

24. How VACCU and OPTIMIZE commands can be used to manage the DELTA lake?

- VACCU command can be used to compact small parquet files, and the OPTIMIZE command can be used to delete
- parquet files that are marked for deletion/unused.
- VACCU command can be used to delete empty/blank parquet files in a delta table.
- OPTIMIZE command can be used to update stale statistics
- on a delta table.
- VACCU command can be used to compress the parquet files to reduce the size of the table, OPTIMIZE command can be used to cache frequently delta tables for better
- performance.
- VACCU command can be used to delete empty/blank parquet files in a delta table, OPTIMIZE command can be
- used to cache frequently delta tables for better performance.
- OPTIMIZE command can be used to compact small parquet files, and the VACCU command can be used to delete parquet files that are marked for deletion/unused. ✓
- (Correct)

CCA-V (9)

CCE-AppDS (1)

CCE-N (1)

CCE-V (5)

CCP-AppDS (2)

CCP-N (4)

CCP-V (6)

CIW (8)

CIW User Interface Designer (1)

CIW Web and Mobile Design Series (4)

CIW Web Development Series (2)

Web Security Series (1)

CloudBees (2)

Jenkins Engineer (2)

Cloudera (1)

Cloudera Certified Associate (CCA) (1)

CMT Association (1)

CMT Program (1)

CNCF (3)

CKA Certification (2)

CompTIA (261)

A+ (53)

CASP (5)

CompTIA Advanced Security Practitioner (2)

CompTIA CASP (16)

Question was not answered

Explanation:

VACCUUM:

You can remove files no longer referenced by a Delta table and are older than the retention threshold by running the vacuum command on the table. vacuum is not triggered automatically. The default retention threshold for the files is 7 days. To change this behavior, see Configure data retention for time travel.

OPTIMIZE:

Using OPTIMIZE you can compact data files on Delta Lake, this can improve the speed of read queries on the table. Too many small files can significantly degrade the performance of the query.

25. Which of the following statements are correct on how Delta Lake implements a lake house?

- Delta lake uses a proprietary format to write data, optimized for cloud storage
- Using Apache Hadoop on cloud object storage
- Delta lake always stores meta data in memory vs storage
- Delta lake uses open source, open format, optimized cloud storage and scalable meta data 
- Delta lake stores data and meta data in computes memory

Question was not answered

Explanation:

Delta lake is

- Open source
- Builds up on standard data format
- Optimized for cloud object storage
- Built for scalable metadata handling Delta lake is not
- Proprietary technology
- Storage format
- Storage medium

CompTIA Cloud Essentials (2)

CompTIA Cloud Essentials+ (2)

CompTIA Cloud+ (12)

CompTIA CloudNetX (1)

CompTIA CySA+ (19)

CompTIA Data+ (5)

CompTIA DataSys+ (1)

CompTIA DataX (1)

CompTIA IT Fundamentals (7)

CompTIA Network+ (36)

CompTIA PenTest+ (17)

CompTIA Project+ (13)

CompTIA Security+ (36)

CompTIA SecurityX Certification (6)

CompTIA Tech+ (1)

CTT+ (1)

Linux+ (14)

Server+ (14)

Confluent (2)

Confluent Certified Administrator (1)

Confluent Certified Developer (1)

Copado (3)

Copado Certifications (1)

Copado Robotic Testing certification (2)

CPA (2)

- Database service or data warehouse

26. What are the different ways you can schedule a job in Databricks workspace?

- Continuous, Incremental
- On-Demand runs, File notification from Cloud object storage
- Cron, On Demand runs 
- Cron, File notification from Cloud object storage
- Once, Continuous

Question was not answered

Explanation:

The answer is, Cron, On-Demand runs

Supports running job immediately or using can be scheduled using CRON syntax

27. Which of the following type of tasks cannot setup through a job?

- Notebook
- DELTA LIVE PIPELINE
- Spark Submit
- Python
- Databricks SQL Dashboard refresh 

Question was not answered

28. Which of the following describes how Databricks Repos can help facilitate CI/CD workflows on the Databricks Lakehouse Platform?

- Databricks Repos can facilitate the pull request, review, and approval process before merging branches

- Databricks Repos can merge changes from a secondary Git branch into a main Git branch

- Databricks Repos can be used to design, develop, and trigger Git automation pipelines

- Databricks Repos can store the single-source-of-truth Git repository

- Databricks Repos can commit or push code changes to trigger a CI/CD process 

Question was not answered

Explanation:

Certified public accountant (2)

CrowdStrike (8)

CCFR (2)

CrowdStrike Falcon
Certification Program (6)

CSA (5)

Cloud Security Knowledge (3)

Zero Trust (2)

CWAP (2)

CWAP (Wi-Fi Analysis) (1)

CWAP Certifications (1)

CWNP (20)

CWDP Certifications (5)

CWICP Certification (1)

CWISA Certification (1)

CWNA Certification (4)

CWS Certification (2)

CWSA Certification (1)

CWSP Certification (3)

CWT Certification (2)

CWTS Certification (1)

Cyber AB (1)

CMMC (1)

CyberArk (24)

CyberArk CDE Recertification
(3)

CyberArk Certified Delivery
Engineer (CDE) (5)

CyberArk Defender (9)

Answer is Databricks Repos can commit or push code changes to trigger a CI/CD process See below diagram to understand the role Databricks Repos and Git provider plays when building a CI/CD workflow.

All the steps highlighted in yellow can be done Databricks Repo, all the steps highlighted in Gray are done in a git provider like Github or Azure Devops.



Diagram

Description automatically generated



**Real Databricks Certified
Data Engineer Associate
Exam Dumps - Complete
the Databricks
Certification Finally**

**Databricks Certified Data
Engineer Professional
Exam Dumps: 278 Practice
Exam Questions and
Answers**



Tags: [Databricks Certified Data Engineer Professional Dumps](#)

Related Posts

Pass Your
Databricks
Certified
Professional
Data Engineer
Exam with
Confidence:
Practice with

Databricks
Certified Data
Engineer
Associate
Dumps
Updated
(V9.02) –
Leverage the

Databricks
Machine
Learning
Associate
Dumps
(V9.02) 2025
– Pass Your
Databricks

- CyberArk Secrets Manager (1)
- CyberArk Sentry (6)
- DAMA (4)
- CDMP Certification (4)
- Dassault Systemes (1)
- CATIA (1)
- Databricks (16)
- Data Analyst (2)
- Databricks Certification (10)
- Generative AI Engineer (1)
- Machine Learning Associate (2)
- ML Data Scientist (1)
- DELL EMC (251)
- Backup Recovery (EMCBA) (1)
- Cloud Architect (EMCCA) (7)
- Data Scientist (1)
- DCA-DPM (2)
- DCA-ISM (15)
- DCA-Networking (2)
- DCE (2)
- DCS-IE (22)
- DCS-SA (4)
- DCS-TA (7)
- DECA-CIS (5)
- DECA-CSHC (2)
- DECA-PowerEdge (3)
- DECE-IE (1)